

Natural Language Processing (almost) from Scratch

Ronan Collobert

RONAN@COLLOBERT.COM

NEC Labs America, Princeton NJ.

Jason Weston

JWESTON@GOOGLE.COM

Google, New York, NY.

Léon Bottou

LEON@BOTTOU.ORG

Michael Karlen

MICHAEL.KARLEN@GMAIL.COM

Koray Kavukcuoglu[†]

KORAY@CS.NYU.EDU

Pavel Kuksa[‡]

PKUKSA@CS.RUTGERS.EDU

NEC Labs America, Princeton NJ.

Editor:

Abstract

We propose a unified neural network architecture and learning algorithm that can be applied to various natural language processing tasks including: part-of-speech tagging, chunking, named entity recognition, and semantic role labeling, achieving or exceeding state-of-the-art performance in each on four benchmark tasks. Our goal was to design a flexible architecture that can *learn* representations useful for the tasks, thus avoiding excessive task-specific feature engineering (and therefore disregarding a lot of prior knowledge). Instead of exploiting man-made input features carefully optimized for each task, our system learns internal representations on the basis of vast amounts of mostly unlabelled training data. This work is then used as a basis for building a freely available tagging system with excellent performance while requiring minimal computational resources.

Keywords: Natural Language Processing, Neural Networks, Deep Learning

1. Introduction

Who nowadays still hopes to design a computer program able to convert a piece of English text into a computer-friendly data structure that unambiguously and completely describes the meaning of the text? Among numerous problems, no consensus has emerged about the form of such a data structure. Until such fundamental problems are resolved, computer scientists must settle for reduced objectives: extracting simpler representations describing restricted aspects of the textual information.

These simpler representations are often motivated by specific applications, for instance, bag-of-words variants for information retrieval. These representations can also be motivated by our belief that they capture something more general about the natural language. They can describe syntactical information (e.g. part-of-speech tagging, chunking, and parsing) or semantic information (e.g. word-sense disambiguation, semantic role labeling, named entity

[†]. Koray Kavukcuoglu is also with New York University, New York, NY.

[‡]. Pavel Kuksa is also with Rutgers University, New Brunswick, NJ.

extraction, and anaphora resolution). Text corpora have been manually annotated with such data structures in order to compare the performance of various systems. The availability of standard benchmarks has stimulated research in Natural Language Processing (NLP).

Many researchers interpret such reduced objectives as stepping stones towards the goal of understanding natural languages. Real-world NLP applications have been approached by smartly reconfiguring these systems and combining their outputs. The benchmark results then tell only a part of the story because they do not measure how effectively these systems can be reconfigured to address real world tasks.

In this paper, we try to excel in multiple benchmark tasks using a *single learning system*. In fact we view the benchmarks as indirect measurements of the relevance of the internal representations discovered by the learning procedure, and we posit that these intermediate representations are more general than any of the benchmark targets.

Many highly engineered NLP systems address the benchmark tasks using linear statistical models applied to task-specific features. In other words, the researchers themselves discover intermediate representations by engineering ad-hoc features. These features are often derived from the output of preexisting systems, leading to complex runtime dependencies. This approach is effective because researchers leverage a large body of linguistic knowledge. On the other hand, there is a great temptation to over-engineer the system to optimize its performance on a particular benchmark at the expense of the broader NLP goals.

In this contribution, we describe a unified NLP system that achieves excellent performance on multiple benchmark tasks by discovering its own internal representations. We have avoided engineering features as much as possible and we have therefore ignored a large body of linguistic knowledge. Instead we reach state-of-the-art performance levels by transferring intermediate representations discovered on massive unlabeled datasets. We call this approach “almost from scratch” to emphasize this reduced (but still important) reliance on a priori NLP knowledge.

The paper is organized as follows. Section 2 describes the benchmark tasks of interest. Section 3 describes the unified model and reports benchmark results obtained with supervised training. Section 4 leverages very large unlabeled datasets to train the model on a language modeling task. Vast performance improvements are demonstrated by transferring the unsupervised internal representations into the supervised benchmark models. Section 5 investigates multitask supervised training. Section 6 evaluates how much further improvement could be achieved by incorporating classical NLP engineering tricks into our systems. We then conclude with a short discussion section.

2. The Benchmark Tasks

In this section, we briefly introduce four classical NLP tasks on which we will benchmark our architectures within this paper: Part-Of-Speech tagging (POS), chunking (CHUNK), Named Entity Recognition (NER) and Semantic Role Labeling (SRL). For each of them, we consider a classical experimental setup and give an overview of state-of-the-art systems on this setup. The experimental setups are summarized in Table 1, while state-of-the-art systems are reported in in Table 2.

Task	Benchmark	Dataset	Training set (#tokens)	Test set (#tokens)	(#tags)
POS	Toutanova et al. (2003)	WSJ	sections 0–18 (912,344)	sections 22–24 (129,654)	(45)
Chunking	CoNLL 2000	WSJ	sections 15–18 (211,727)	section 20 (47,377)	(42) (IOBES)
NER	CoNLL 2003	Reuters	“eng.train” (203,621)	“eng.testb” (46,435)	(17) (IOBES)
SRL	CoNLL 2005	WSJ	sections 2–21 (950,028)	section 23 + 3 Brown sections (63,843)	(186) (IOBES)

Table 1: Experimental setup: for each task, we report the standard benchmark we used, the dataset it relates to, as well as training and test information.

System	Accuracy
Shen et al. (2007)	97.33%
Toutanova et al. (2003)	97.24%
Giménez and Màrquez (2004)	97.16%
(a) POS	
System	F1
Shen and Sarkar (2005)	95.23%
Sha and Pereira (2003)	94.29%
Kudo and Matsumoto (2001)	93.91%
(b) CHUNK	
System	F1
Ando and Zhang (2005)	89.31%
Florian et al. (2003)	88.76%
Kudo and Matsumoto (2001)	88.31%
(c) NER	
System	F1
Koomen et al. (2005)	77.92%
Pradhan et al. (2005)	77.30%
Haghighi et al. (2005)	77.04%
(d) SRL	

Table 2: State-of-the-art systems on four NLP tasks. Performance is reported in per-word accuracy for POS, and F1 score for CHUNK, NER and SRL. Systems in bold will be referred as *benchmark systems* in the rest of the paper (see text).

2.1 Part-Of-Speech Tagging

POS aims at labeling each word with a unique tag that indicates its *syntactic role*, e.g. plural noun, adverb, ... A classical benchmark setup is described in detail by Toutanova

et al. (2003). Sections 0–18 of WSJ data are used for training, while sections 19–21 are for validation and sections 22–24 for testing.

The best POS classifiers are based on classifiers trained on windows of text, which are then fed to a bidirectional decoding algorithm during inference. Features include preceding and following tag context as well as multiple words (bigrams, trigrams...) context, and handcrafted features to deal with unknown words. Toutanova et al. (2003), who use maximum entropy classifiers, and a bidirectional dependency network (Heckerman et al., 2001) at inference, reach 97.24% per-word accuracy. Giménez and Màrquez (2004) proposed a SVM approach also trained on text windows, with bidirectional inference achieved with two Viterbi decoders (left-to-right and right-to-left). They obtained 97.16% per-word accuracy. More recently, Shen et al. (2007) pushed the state-of-the-art up to 97.33%, with a new learning algorithm they call *guided learning*, also for bidirectional sequence classification.

2.2 Chunking

Also called shallow parsing, chunking aims at labeling segments of a sentence with syntactic constituents such as noun or verb phrase (NP or VP). Each word is assigned only one unique tag, often encoded as a begin-chunk (e.g. B-NP) or inside-chunk tag (e.g. I-NP). Chunking is often evaluated using the CoNLL 2000 shared task¹. Sections 15–18 of WSJ data are used for training and section 20 for testing. Validation is achieved by splitting the training set.

Kudoh and Matsumoto (2000) won the CoNLL 2000 challenge on chunking with a F1-score of 93.48%. Their system was based on Support Vector Machines (SVMs). Each SVM was trained in a pairwise classification manner, and fed with a window around the word of interest containing POS and words as features, as well as surrounding tags. They perform dynamic programming at test time. Later, they improved their results up to 93.91% (Kudo and Matsumoto, 2001) using an ensemble of classifiers trained with different tagging conventions Section 3.2.3.

Since then, a certain number of systems based on second-order random fields were reported (Sha and Pereira, 2003; McDonald et al., 2005; Sun et al., 2008), all reporting around 94.3% F1 score. These systems use features composed of words, POS tags, and tags.

More recently, Shen and Sarkar (2005) obtained 95.23% using a voting classifier scheme, where each classifier is trained on different tag representations (IOB1, IOB2, ...). They use POS features coming from an external tagger, as well carefully hand-crafted *specialization* features which again change the data representation by concatenating some (carefully chosen) chunk tags or some words with their POS representation. They then build trigrams over these features, which are finally passed through a Viterbi decoder a test time.

2.3 Named Entity Recognition

NER labels atomic elements in the sentence into categories such as “PERSON”, “COMPANY”, or “LOCATION”. As in the chunking task, each word is assigned a tag prefixed

1. See <http://www.cnts.ua.ac.be/conll2000/chunking>.

by an indicator of the beginning or the inside of an entity. The CoNLL 2003 setup² is a NER benchmark dataset based on Reuters data. The contest provides training, validation and testing sets.

Florian et al. (2003) presented the best system at the NER CoNLL 2003 challenge, with 88.76% F1 score. They used a combination of various machine-learning classifiers. Features they picked included words, POS tags, CHUNK tags, prefixes and suffixes, a large gazetteer (not provided by the challenge), as well as the output of two other NER classifiers trained on richer datasets. Chieu (2003), the second best performer of CoNLL 2003 (88.31% F1), also used an external gazetteer (their performance goes down to 86.84% with no gazetteer) and a several hand-chosen features.

Later, Ando and Zhang (2005) reached 89.31% F1 with a semi-supervised approach. They trained jointly a linear model on NER with a linear model on two auxiliary unsupervised tasks. They also performed Viterbi decoding at test time. The unlabeled corpus was 27M words taken from Reuters. Features included words, POS tags, suffixes and prefixes or CHUNK tags, but overall were less specialized than CoNLL 2003 challengers.

2.4 Semantic Role Labeling

SRL aims at giving a semantic role to a syntactic constituent of a sentence. In the PropBank (Palmer et al., 2005) formalism one assigns roles ARG0-5 to words that are arguments of a verb (or more technically, a *predicate*) in the sentence, e.g. the following sentence might be tagged “[John]_{ARG0} [ate]_{REL} [the apple]_{ARG1} ”, where “ate” is the predicate. The precise arguments depend on a verb’s *frame* and if there are multiple verbs in a sentence some words might have multiple tags. In addition to the ARG0-5 tags, there are several modifier tags such as ARG-M-LOC (locational) and ARG-M-TMP (temporal) that operate in a similar way for all verbs. We picked CoNLL 2005³ as our SRL benchmark. It takes sections 2–21 of WSJ data as training set, and section 24 as validation set. A test set composed of section 23 of WSJ concatenated with 3 sections from the Brown corpus is also provided by the challenge.

State-of-the-art SRL systems consist of several stages: producing a parse tree, identifying nodes of the parse tree which are arguments for a given verb chosen beforehand, and finally classifying these arguments into the right SRL tags. Existing SRL systems use a *lot* of features, including POS tags, head words, phrase type, path in the parse tree from the verb to the node.... Koomen et al. (2005) hold the state-of-the-art with Winnow-like (Littlestone, 1988) classifiers, followed by a decoding stage based on an integer program that enforces specific constraints on SRL tags. They reach 77.92% F1 on CoNLL 2005, thanks to the five top parse trees produced by the Charniak (2000) parser (only the first one was provided by the contest) as well as the Collins (1999) parse tree. Pradhan et al. (2005) obtained 77.30% F1 with a system based on SVM classifiers, and the two parse trees provided for the SRL task. In the same spirit, the system proposed by Haghighi et al. (2005) was based on log-linear models on each tree node, re-ranked globally with a dynamic algorithm. It reached 77.04% using again the five top Charniak parse trees.

2. See <http://www.cnts.ua.ac.be/conll2003/ner>.

3. See <http://www.lsi.upc.edu/~srlconll>.

2.5 Discussion

When participating to an (open) challenge, it is legitimate to increase generalization by all means. It is thus not surprising to see many top CoNLL systems using *external labeled data*, like additional NER classifiers for the NER architecture of Florian et al. (2003) or additional parse trees for SRL systems (Koomen et al., 2005). Combining multiple systems or tweaking carefully features is also a common approach, like in the chunking top system (Shen and Sarkar, 2005).

However, when *comparing* systems, we do not learn anything of the quality of each system if they were trained with *different* labeled data or if some were engineered to death. For that reason, we will refer to *benchmark systems*, that is, top existing systems which avoid usage of external data and have been *well-established* in the NLP field: (Toutanova et al., 2003) for POS and (Sha and Pereira, 2003) for chunking. For NER we consider (Ando and Zhang, 2005) as they were using only *unlabeled* data. We picked (Koomen et al., 2005) for SRL, keeping in mind they use 4 additional parse trees not provided by the challenge. These benchmark systems will serve us as baseline references in our experiments. We marked them in bold in Table 2.

3. The Networks

All the NLP tasks above can be seen as tasks assigning labels to words. The traditional NLP approach is: extract from the sentence a rich set of hand-designed features which are then fed to a classical *shallow* classification algorithm, e.g. a Support Vector Machine (SVM), often with a linear kernel. The choice of features is a completely empirical process, mainly based first on linguistic intuition, and then trial and error, and the feature selection is task dependent, implying additional research for each new NLP task. Complex tasks like SRL then require a large number of possibly complex features (e.g., extracted from a parse tree) which makes such systems slow and intractable for large-scale applications.

Instead, we advocate a radically different approach: as input we will try to pre-process our features as little as possible and then use a deep neural network (NN) architecture, trained in an end-to-end fashion. The architecture takes the input sentence and learns several layers of feature extraction that process the inputs. The features in the deep layers of the network are *automatically trained* by backpropagation to be relevant to the task. We describe in this section a general deep architecture suitable for all our NLP tasks, which is generalizable to other NLP tasks as well.

Our architecture is summarized in Figure 1 and Figure 2. The first layer extracts features for each word. The second layer extracts features from a window of words or from the whole sentence, treating it as a *sequence* with local and global structure (i.e., it is not treated like a bag of words). The following layers are classical NN layers.

3.1 Transforming Words into Feature Vectors

One of the essential keypoints of our architecture is its ability to perform well with the sole use of raw words, without any engineered features. The ability for our method to learn good word representations is thus crucial to our approach. For efficiency, words are fed to our architecture as indices taken from a finite dictionary \mathcal{D} . Obviously, a simple index does

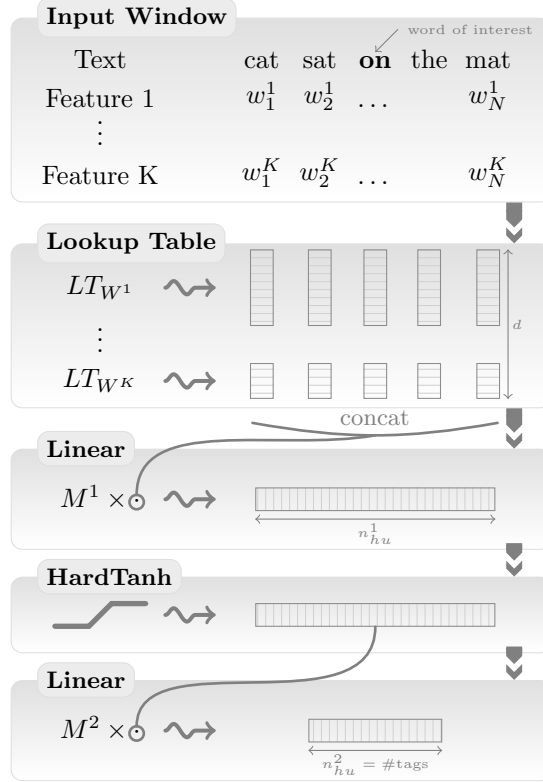


Figure 1: Window approach network.

not carry much useful information about the word. However, the first layer of our network maps each of these word indices into a feature vector, by a lookup table operation. Given a task of interest, a *relevant* representation of each word is then given by the corresponding lookup table feature vector, which is *trained* by backpropagation.

More formally, for each word $w \in \mathcal{D}$, an internal d_{word} -dimensional feature vector representation is given by the *lookup table* layer $LT_W(\cdot)$:

$$LT_W(w) = W_{\cdot w},$$

where $W \in \mathbb{R}^{d_{word} \times |\mathcal{D}|}$ is a matrix of parameters to be *learned*, $W_{\cdot w} \in \mathbb{R}^{d_{word}}$ is the w^{th} column of W and d_{word} is the word vector size (a hyper-parameter to be chosen by the user). Given a sentence or any sequence of T words $[w]_1^T$ in \mathcal{D} , the lookup table layer applies the same operation for each word in the sequence, producing the following output matrix:

$$LT_W([w]_1^T) = \begin{pmatrix} W_{\cdot [w]_1} & W_{\cdot [w]_2} & \dots & W_{\cdot [w]_T} \end{pmatrix}. \quad (1)$$

This matrix can be then fed to further neural network layers, as we will see below.

3.1.1 EXTENDING TO ANY DISCRETE FEATURES

One might want to provide features other than words if these features are suspected to be helpful for the task of interest. For example, for the NER task, one could provide a feature




Figure 2: Sentence approach network.

which says if a word is in a gazetteer or not. Another common practice is to introduce some basic pre-processing, such as word-stemming or dealing with upper and lower case. In this latter option, the word would be then represented by three discrete features: its lower case stemmed root, its lower case ending, and a capitalization feature.

Generally speaking, we can consider a word as represented by K discrete features $\mathbf{w} \in \mathcal{D}^1 \times \dots \times \mathcal{D}^K$, where \mathcal{D}^k is the dictionary for the k^{th} feature. We associate to each feature a lookup table $LT_{W^k}(\cdot)$, with parameters $W^k \in \mathbb{R}^{d_{wrd}^k \times |\mathcal{D}^k|}$ where $d_{wrd}^k \in \mathbb{N}$ is a user-specified

vector size. Given a word \mathbf{w} , a feature vector of dimension $d_{wrd} = \sum_k d_{wrd}^k$ is then obtained by concatenating all lookup table outputs:

$$LT_{W^1, \dots, W^K}(\mathbf{w}) = \begin{pmatrix} LT_{W^1}(w_1) \\ \vdots \\ LT_{W^K}(w_K) \end{pmatrix} = \begin{pmatrix} W_{w_1}^1 \\ \vdots \\ W_{w_K}^K \end{pmatrix}.$$

The matrix output of the lookup table layer for a sequence of words $[\mathbf{w}]_1^T$ is then similar to (1), but where extra rows have been added for each discrete feature:

$$LT_{W^1, \dots, W^K}([\mathbf{w}]_1^T) = \begin{pmatrix} W_{[w_1]_1}^1 & \dots & W_{[w_1]_T}^1 \\ \vdots & & \vdots \\ W_{[w_K]_1}^K & \dots & W_{[w_K]_T}^K \end{pmatrix}. \quad (2)$$

These vector features learnt by the lookup table learn features for words. Now, we want to use these trainable features as input to further layers of trainable feature extractors, that can represent groups of words and then finally sentences.

3.2 Extracting Higher Level Features from Word Feature Vectors

Feature vectors produced by the lookup table layer need to be combined in subsequent layers of the neural network to produce a tag decision for each word in the sentence. Producing tags for each element in variable length sequences (here, a sentence is a sequence of words) is a classical problem in machine-learning. We consider two common approaches which tag *one word at the time*: a window approach, and a (convolutional) sentence approach.

3.2.1 WINDOW APPROACH

A window approach assumes the tag of a word depends mainly on its neighboring words. Given a word to tag, we consider a fixed-size k_{sz} (hyper-parameter) window of words around this word. Each word in the window is first passed through the lookup table layer (1) or (2), producing a matrix of word features of fixed size $d_{wrd} \times k_{sz}$. This matrix can be viewed as a $d_{wrd} k_{sz}$ -dimensional vector by concatenating each column vector, which can be fed to further neural network layers.

More formally, we denote \mathbf{z}^l to be the output of the l^{th} layer of our neural network and adopt the notation $\langle \cdot \rangle_t^{d_{win}}$ for designating the d_{win} concatenated column vectors of a given matrix around the t^{th} column vector. Then, considering the t^{th} word to tag, the word feature window is:

$$\mathbf{z}^1 = \langle LT_W([\mathbf{w}]_1^T) \rangle_t^{d_{win}} = \begin{pmatrix} W_{[w]_{t-d_{win}/2}} \\ \vdots \\ W_{[w]_t} \\ \vdots \\ W_{[w]_{t+d_{win}/2}} \end{pmatrix}. \quad (3)$$

The fixed size vector \mathbf{z}^1 can be fed to one or several classical neural network layers which perform affine transformations over their inputs:

$$\mathbf{z}^l = M^l \mathbf{z}^{l-1} + \mathbf{b}^l, \quad (4)$$

where $M^l \in \mathbb{R}^{n_{hu}^l \times n_{hu}^{l-1}}$ and $\mathbf{b}^l \in \mathbb{R}^{n_{hu}^l}$ are the parameters to be *trained*. The hyper-parameter n_{hu}^l is usually called the *number of hidden units* of the l^{th} layer.

Several of these layers are often stacked, to extract highly non-linear features. If the l^{th} classical layer is not the last layer, a non-linearity function must follow, or our network would be a simple linear model. We chose a “hard” version of the hyperbolic tangent version, which has the advantage of being slightly cheaper to compute while leaving the generalization performance unchanged.

$$\text{HardTanh}(x) = \begin{cases} -1 & \text{if } x < -1 \\ x & \text{if } -1 \leq x \leq 1 \\ 1 & \text{if } x > 1 \end{cases} . \quad (5)$$

Finally, the output size of the last layer L of our network is equal to the number of possible tags for the task of interest. Each output can be then interpreted as a *score* of the corresponding tag (given the input of the network), thanks to a carefully chosen cost function that we will describe later in this section.

Remark 1 (Border Effects) *The feature window (3) is not well defined for words near the beginning or the end of a sentence. To circumvent this problem, we augment the sentence with a special “PADDING” word replicated $d_{win}/2$ times at the beginning and the end.*

3.2.2 SENTENCE APPROACH

We will see in the experimental section that a window approach performs well for most natural language processing tasks we are interested in. However this approach fails with SRL, where the tag of a word depends on a verb (or, more correctly, predicate) chosen beforehand in the sentence: if the verb falls outside the window, one cannot expect this word to be tagged correctly. In this particular case, tagging a word requires the consideration of the *whole* sentence. When using neural networks, the natural choice to tackle this problem becomes a convolutional approach, first introduced by Waibel et al. (1989) and also called Time Delay Neural Networks (TDNNs) in the literature.

We describe in details our convolutional network below. It successively takes the complete sentence, passes it through the lookup table layer (1), produces local features around each word of the sentence thanks to convolutional layers, combines these feature into a global feature vector which can be then fed to classical affine layers (4). In the semantic role labeling case, this operation is performed for each word in the sentence, and for each verb in the sentence. It is thus necessary to encode in the network architecture which verb we are considering in the sentence, and which word we want to tag. For that purpose, each word at position i in the sentence is augmented with two features in the way described in Section 3.1.1. These features encode the relative distances $i - pos_v$ and $i - pos_w$ with respect to the chosen verb at position pos_v , and the word to tag at position pos_w respectively.

A convolutional layer can be seen as a generalization of a window approach: given a sequence represented by columns in a matrix Z^{l-1} (in our the lookup table matrix (1)), a matrix-vector operation as in (4) is applied to each window of successive windows in the sequence. Using previous notations, the t^{th} output column of the l^{th} layer can be computed

as:

$$Z_{.t}^l = M^l \langle Z^{l-1} \rangle_t^{d_{win}} + \mathbf{b}^l, \quad (6)$$

where the weight matrix M^l is the same across all windows t in the sequence. Convolutional layers extract local features around each window of the given sequence. As for classical affine layers (4), convolutional layers are often stacked to extract higher level features. In this case, each layer must be followed by a non-linearity (5) or the network would be equivalent to one convolutional layer.

The size of the output (6) depends on the number of words in the sentence fed to the network. Local feature vectors extracted by the convolutional layers have to be combined to obtain a global feature vector, with a fixed-size independent of the sentence length, in order to apply subsequent classical affine layers. Traditional convolutional networks often apply an average (possibly weighted) or a max operation over the “time” t of the sequence (6). (Here, “time” just means the position in the sentence, this term stems from the use of convolutional layers in e.g. speech data where the sequence occurs over time.) The average operation does not make much sense in our case, as in general most words in the sentence do not have any influence on the semantic role of a given word to tag. Instead, we used a max approach, which forces the network to capture the most useful local features for the task at hand that are produced by the convolutional layers:

$$z_i^l = \max_t Z_{i,t}^{l-1} \quad 1 \leq i \leq n_{hu}^{l-1}. \quad (7)$$

This fixed-sized global feature vector can be then fed to classical affine network layers (4). As in the window approach, we then finally produce one score per possible tag for the given task.

Remark 2 *The same border effects arise in the convolution operation (6) as in the window approach (3). We again work around this problem by padding the sentences with a special word.*

3.2.3 TAGGING SCHEMES

As explained earlier, the network output layers compute scores for all the possible tags for the task of interest. In the window approach, these tags apply to the word located in the center of the window. In the (convolutional) sentence approach, these tags apply to the word designated by additional markers in the network input.

The POS task consists indeed of marking the syntactic role of each word. However, the remaining three tasks associate tags with segments of a sentence. This is usually achieved by using special tagging schemes to identify the segment boundaries. Several such schemes have been defined (IOB, IOB2, IOE, IOBES, ...) without clear conclusion as to which scheme is better in general. State-of-the-art performance is sometimes obtained by combining classifiers trained with different tagging schemes (e.g. Kudo and Matsumoto, 2001).

The ground truth for the NER, CHUNK, and SRL tasks is provided using two different tagging schemes. In order to eliminate this additional source of variations, we have decided to use the most expressive IOBES tagging scheme for all tasks. For instance, in the CHUNK task, we describe noun phrases using four different tags. Tag “S-NP” is used to mark a noun phrase containing a single word. Otherwise tags “B-NP”, “I-NP”, and “E-NP” are used

to mark the first, intermediate and last words of the noun phrase. An additional tag “O” marks words that are not members of a chunk. During testing, these tags are then converted to the original IOB tagging scheme and fed to the standard performance evaluation script.

3.3 Training

All our neural networks are trained by maximizing a likelihood over the training data, using stochastic gradient ascent. If we denote θ to be all the trainable parameters of the network, which are trained using a training set \mathcal{T} we want to maximize:

$$\theta \mapsto \sum_{(\mathbf{x}, y) \in \mathcal{T}} \log p(y | \mathbf{x}, \theta), \quad (8)$$

where \mathbf{x} corresponds to either a training word window or a sentence and its associated features, and y represents the corresponding tag. The probability $p(\cdot)$ is computed from the outputs of the neural network. We will see in this section two ways of interpreting neural network outputs as probabilities.

3.3.1 ISOLATED TAG CRITERION

In this approach, each word in a sentence is considered independently. Given an input example \mathbf{x} , the network with parameters θ outputs a score that we denote $f(\mathbf{x}, i, \theta)$, for the i^{th} tag with respect to the task of interest. This score can be interpreted as a conditional tag probability $p(i | \mathbf{x}, \theta)$ by applying a softmax (Bridle, 1990) operation over all the tags:

$$p(i | \mathbf{x}, \theta) = \frac{e^{f(\mathbf{x}, i, \theta)}}{\sum_j e^{f(\mathbf{x}, j, \theta)}}. \quad (9)$$

Defining the log-add operation as

$$\text{logadd}_i z_i = \log\left(\sum_i e^{z_i}\right), \quad (10)$$

we can express the log-likelihood for one training example (\mathbf{x}, y) as follows:

$$\log p(y | \mathbf{x}, \theta) = f(\mathbf{x}, y, \theta) - \text{logadd}_j f(\mathbf{x}, j, \theta). \quad (11)$$

While this training criterion, often referred as *cross-entropy* is widely used for classification problems, it might not be ideal in our case, where often there is a correlation between the tag of a word in a sentence and its neighboring tags. We now describe another common approach for neural networks which enforces dependencies between the predicted tags in a sentence.

3.3.2 SENTENCE TAG CRITERION

In tasks like chunking, NER or SRL we know that there are dependencies between word tags in a sentence: not only are tags organized in chunks, but some tags cannot follow other tags. Training using an isolated tag approach discards this kind of labeling information. We

consider a training scheme which takes in account the sentence structure: given predictions for *all tags* by our network for *all words* in a sentence, and given a score for going from one tag to another tag, we want to encourage valid paths of tags during training, while discouraging all other paths.

We define $f([\mathbf{x}]_1^T, i, t, \boldsymbol{\theta})$, the score output by the network with parameters $\boldsymbol{\theta}$, for the sentence $[\mathbf{x}]_1^T$ and for the i^{th} tag, at the t^{th} word. We introduce a transition score A_{ij} for jumping from i to j tags in successive words, and an initial score A_{i0} for starting from the i^{th} tag. As the transition scores are going to be trained (as are all network parameters $\boldsymbol{\theta}$), we define $\tilde{\boldsymbol{\theta}} = \boldsymbol{\theta} \cup \{A_{ij} \forall i, j\}$. The score of a sentence $[\mathbf{x}]_1^T$ along a path of tags $[j]_1^T$ is then given by the sum of transition scores and network scores:

$$s([\mathbf{x}]_1^T, [j]_1^T, \tilde{\boldsymbol{\theta}}) = \sum_{t=1}^T \left(A_{[j]_{t-1}[j]_t} + f([\mathbf{x}]_1^T, [j]_t, t, \boldsymbol{\theta}) \right). \quad (12)$$

Exactly as for the isolated tag criterion (11) where we were normalizing with respect to all *tags* using a softmax (9), this score can be interpreted as a conditional *tag path* probability, by applying a softmax over all possible *tag paths* $[j]_1^T$. Taking the log, the conditional probability of the true path $[y]_1^T$ is given by:

$$\log p([y]_1^T | [\mathbf{x}]_1^T, \tilde{\boldsymbol{\theta}}) = s([\mathbf{x}]_1^T, [y]_1^T, \tilde{\boldsymbol{\theta}}) - \underset{\forall [j]_1^T}{\text{logadd}} s([\mathbf{x}]_1^T, [j]_1^T, \tilde{\boldsymbol{\theta}}). \quad (13)$$

While the number of terms in the logadd operation (11) was equal to the number of tags, it grows exponentially with the length of the sentence in (13). Fortunately, one can compute it in linear time with the following classical recursion over t , taking advantage of the associativity and distributivity on the semi-ring⁴ $(\mathbb{R} \cup \{-\infty\}, \text{logadd}, +)$:

$$\begin{aligned} \delta_t(k) &\triangleq \underset{\{[j]_1^t \cap [j]_t = k\}}{\text{logadd}} s([\mathbf{x}]_1^t, [j]_1^t, \tilde{\boldsymbol{\theta}}) \\ &= \underset{i}{\text{logadd}} \underset{\{[j]_1^t \cap [j]_{t-1} = i \cap [j]_t = k\}}{\text{logadd}} s([\mathbf{x}]_1^t, [j]_1^{t-1}, \tilde{\boldsymbol{\theta}}) + A_{[j]_{t-1}k} + f([\mathbf{x}]_1^t, k, t, \boldsymbol{\theta}) \\ &= \underset{i}{\text{logadd}} \delta_{t-1}(i) + A_{ik} + f([\mathbf{x}]_1^t, k, t, \boldsymbol{\theta}) \\ &= f([\mathbf{x}]_1^t, k, t, \boldsymbol{\theta}) + \underset{i}{\text{logadd}} (\delta_{t-1}(i) + A_{ik}) \quad \forall k, \end{aligned} \quad (14)$$

followed by the termination

$$\underset{\forall [j]_1^T}{\text{logadd}} s([\mathbf{x}]_1^T, [j]_1^T, \tilde{\boldsymbol{\theta}}) = \underset{i}{\text{logadd}} \delta_T(i). \quad (15)$$

We can now maximize in (8) the log-likelihood (13) over all the training pairs $([\mathbf{x}]_1^T, [y]_1^T)$.

At inference time, given a sentence $[\mathbf{x}]_1^T$ to tag, we have to find the best tag path which minimizes the sentence score (12). In other word, we must find

$$\underset{[j]_1^T}{\text{argmax}} s([\mathbf{x}]_1^T, [j]_1^T, \tilde{\boldsymbol{\theta}}). \quad (16)$$

4. In other words, read logadd as \oplus and $+$ as \otimes .

The Viterbi algorithm is the natural choice for this inference. It corresponds to performing the recursion (14) and (15), but where the logadd is replaced by a max, and then tracking back the optimal path through each max.

Remark 3 (Graph Transformers) *Our approach is a particular case of the discriminative forward training for graph transformers (Bottou et al., 1997; Le Cun et al., 1998). The log-likelihood (13) can be viewed as the difference between the forward score constrained over the valid paths (in our case there is only the labeled path) and the unconstrained forward score (15).*

Remark 4 (Conditional Random Fields) *Conditional Random Fields (CRFs) (Lafferty et al., 2001) over temporal sequences maximize the same likelihood as (13). The CRF model is however linear (which would correspond in our case to a linear neural network). CRFs have been widely used in the NLP world, such as for POS tagging (Lafferty et al., 2001), chunking (Sha and Pereira, 2003), NER (McCallum and Li, 2003) or SRL (Cohn and Blunsom, 2005). Compared to CRFs approaches, we take advantage of our deep network to train appropriate features for each task of interest.*

3.3.3 STOCHASTIC GRADIENT

Maximizing (8) with stochastic gradient (Bottou, 1991) is achieved by iteratively selecting a random example (\mathbf{x}, y) and making a gradient step:

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \lambda \frac{\partial \log p(y | \mathbf{x}, \boldsymbol{\theta})}{\partial \boldsymbol{\theta}}, \quad (17)$$

where λ is a chosen learning rate. Our neural networks described in Figure 1 and Figure 2 are a succession of layers that correspond to successive composition of functions. The neural network is finally composed with the isolated tag criterion (11), or successively composed in the recursion (14) if using the sentence criterion (13). Thus, an *analytical* formulation of the derivative (17) can be computed, by applying the differentiation chain rule through the network, and through the isolated word criterion (11) or through the recurrence (14).

Remark 5 (Differentiability) *Our cost functions are differentiable almost everywhere. Nondifferentiable points arise because we use a “hard” transfer function (5) and because we use a “max” layer (7) in the sentence approach network. Fortunately, stochastic gradient still converges to a meaningful local minimum despite such minor differentiability problems (Bottou, 1991).*

Remark 6 (Modular Approach) *The well known “back-propagation” algorithm (LeCun, 1985; Rumelhart et al., 1986) computes gradients using the chain rule. The chain rule can also be used in a modular implementation.⁵ Our modules correspond to the boxes in Figure 1 and Figure 2. Given derivatives with respect to its outputs, each module can independently compute derivatives with respect to its inputs and with respect to its trainable parameters, as proposed by Bottou and Gallinari (1991). This allows us to easily build variants of our networks.*

5. See <http://torch5.sf.net>.

Approach	POS (PWA)	Chunking (F1)	NER (F1)	SRL (F1)
Benchmark Systems	97.24	94.29	89.31	77.92
NN+ITC	96.31	89.13	79.53	55.40
NN+STC	96.37	90.33	81.47	70.99

Table 3: Comparison in generalization performance of benchmark NLP systems with our neural network (NN) approach, on POS, chunking, NER and SRL tasks. We report results with both the isolated tag criterion (ITC) and the sentence tag criterion (STC). Generalization performance is reported in per-word accuracy rate (PWA) for POS and F1 score for other tasks.

Task	Window/Conv. size	Word dim.	Caps dim.	Hidden units	Learning rate
POS	$d_{win} = 5$	$d^0 = 50$	$d^1 = 5$	$n_{hu}^1 = 300$	$\lambda = 0.01$
CHUNK	”	”	”	”	”
NER	”	”	”	”	”
SRL	”	”	”	$n_{hu}^1 = 300$ $n_{hu}^2 = 500$	”

Table 4: Hyper-parameters of our networks. We report for each task the window size (or convolution size), word feature dimension, capital feature dimension, number of hidden units and learning rate.

Remark 7 (Tricks) *Many tricks have been reported for training neural networks (LeCun et al., 1998), and it is often confusing which one to choose. We employed only two of them: (1) the initialization of the parameters was done according to the fan-in, and (2) the learning rate was divided by the fan-in (Plaut and Hinton, 1987), but stayed fixed during the training.*

3.4 Supervised Benchmark Results

For POS, chunking and NER tasks, we report results with the window architecture described in Section 3.2.1. The SRL task was trained using the sentence approach (Section 3.2.2). Results are reported in Table 3, in per-word accuracy (PWA) for POS, and F1 score for all the other tasks. We performed experiments both with the isolated tag criterion (ITC) and with the sentence tag criterion (STC). The hyper-parameters of our networks are reported in Table 4. All our networks were fed with two raw text features: lower case words, and a capital letter feature. We chose to consider lower case words to limit the number of words in the dictionary. However, to keep some upper case information lost by this transformation, we added a “caps” feature which tells if each word was in low caps, was all caps, had first letter capital, or had one capital. Additionally, all occurrences of sequences of numbers within a word are replaced with the string “NUMBER”, so for example both the words “PS1” and “PS2” would map to the single word “psNUMBER”. We used a dictionary

FRANCE	JESUS	XBOX	REDDISH	SCRATCHED	MEGABITS
454	1973	6909	11724	29869	87025
PERSUADE	THICKETS	DECADENT	WIDESCREEN	ODD	PPA
FAW	SAVARY	DIVO	ANTICA	ANCHIETA	UDDIN
BLACKSTOCK	SYMPATHETIC	VERUS	SHABBY	EMIGRATION	BIOLOGICALLY
GIORGI	JFK	OXIDE	AWE	MARKING	KAYAK
SHAHEED	KHWARAZM	URBINA	THUD	HEUER	MCLARENS
RUMELIA	STATIONERY	EPOS	OCCUPANT	SAMBHAJI	GLADWIN
PLANUM	ILIAS	EGLINTON	REVISED	WORSHIPPERS	CENTRALLY
GOA’ULD	GSNUMBER	EDGING	LEAVENED	RITSUKO	INDONESIA
COLLATION	OPERATOR	FRG	PANDIONIDAE	LIFELESS	MONEO
BACHA	W.J.	NAMSOS	SHIRT	MAHAN	NILGIRIS

Table 5: Word embeddings in the word lookup table of a SRL neural network trained from scratch, with a dictionary of size 100,000. For each column the queried word is followed by its index in the dictionary (higher means more rare) and its 10 nearest neighbors (arbitrary using the Euclidean metric).

containing the 100,000 most common words in WSJ (case insensitive). Words outside this dictionary were replaced by a single special “RARE” word.

Results show that neural networks “out-of-the-box” are behind baseline benchmark systems. Looking at all submitted systems reported on each CoNLL challenge website showed us our networks performance are nevertheless in the performance ballpark of existing approaches. The training criterion which takes into account the sentence structure seems to always boost the performance for all tasks.

The capacity of our network architectures lies mainly in the word lookup table, which contains $50 \times 100,000$ parameters to train. In the WSJ data, 15% of the most common words appear about 90% of the time. Many words appear only a few times. It is thus hopeless to train properly their corresponding 50 dimensional feature vectors in the lookup table. Ideally, we would like semantically similar words to be close in the embedding space represented by the word lookup table: by continuity of the neural network function, tags produced on semantically similar sentences would be identical. We show in Table 5 that it is not the case: neighboring words in the embedding space do not seem to be semantically related.

We will focus in the next section on improving these word embeddings by leveraging unlabeled data. We will see our approach results in a performance boost for all tasks.

Remark 8 (Architectures) *We have only tried a few different architectures by validation. In practice, the choice of hyperparameters such as the number of hidden units, provided they are large enough, has a limited impact on the generalization performance.*

4. Lots of Unlabeled Data

We would like to obtain word embeddings carrying more syntactic and semantic information than shown in Table 5. Since most of the trainable parameters of our system are associated with the word embeddings, these poor results suggest that we should use considerably more training data. Following our NLP *from scratch* philosophy, we now describe how to dramatically improve these embeddings using *very large unlabeled datasets*. We then use these improved embeddings to initialize the word lookup tables of the networks described in Section 3.4.

4.1 Datasets

Our first English corpus is the entire English Wikipedia.⁶ We have removed all paragraphs containing non-roman characters and all MediaWiki markups. The resulting text was tokenized using the Penn Treebank tokenizer script.⁷ The resulting dataset contains about 631 million words. As in our previous experiments, we use a dictionary containing the 100,000 most common words in WSJ, with the same processing of capitals and numbers. Again, words outside the dictionary were replaced by the special “RARE” word.

Our second English corpus is composed by adding an extra 221 million words extracted from the Reuters RCV1 (Lewis et al., 2004) dataset.⁸ We also extended the dictionary to 130,000 words by adding the 30,000 most common words in Reuters. This is useful in order to determine whether improvements can be achieved by further increasing the unlabeled dataset size.

4.2 Ranking Criterion versus Entropy Criterion

We used these unlabeled datasets to train *language models* that compute *scores* describing the acceptability of a piece of text. These language models are again large neural networks using the window approach described in Section 3.2.1 and in Figure 1. As in the previous section, most of the trainable parameters are located in the lookup tables.

Very similar language models were already proposed by Bengio and Ducharme (2001) and Schwenk and Gauvain (2002). Their goal was to estimate the *probability* of a word given the previous words in a sentence. Estimating conditional probabilities suggests a cross-entropy criterion similar to those described in Section 3.3.1. Because the dictionary size is large, computing the normalization term can be extremely demanding, and sophisticated approximations are required. More importantly for us, neither work leads to significant word embeddings being reported.

Shannon (1951) has estimated the entropy of the English language between 0.6 and 1.3 bits per character by asking human subjects to guess upcoming characters. Cover and King (1978) give a lower bound of 1.25 bits per character using a subtle gambling approach. Meanwhile, using a simple word trigram model, Brown et al. (1992) reach 1.75 bits per character. Teahan and Cleary (1996) obtain entropies as low as 1.46 bits per character using variable length character *n*-grams. The human subjects rely of course on all their

6. Available at <http://download.wikimedia.org>.

7. Available at <http://www.cis.upenn.edu/~treebank/tokenization.html>.

8. Now available at <http://trec.nist.gov/data/reuters/reuters.html>.

knowledge of the language and of the world. Can we learn the grammatical structure of the English language and the nature of the world by leveraging the 0.2 bits per character that separate human subjects from simple n-gram models? Since such tasks certainly require high capacity models, obtaining sufficiently small confidence intervals on the test set entropy may require prohibitively large training sets.⁹ The entropy criterion lacks dynamical range because its numerical value is largely determined by the most frequent phrases. In order to learn syntax, rare but legal phrases are no less significant than common phrases.

It is therefore desirable to define alternative training criteria. We propose here to use a *pairwise ranking* approach (Cohen et al., 1998). We seek a network that computes a higher score when given a legal phrase than when given an incorrect phrase. Because the ranking literature often deals with information retrieval applications, many authors define complex ranking criteria that give more weight to the ordering of the best ranking instances (see Burges et al., 2007; Cl  men  on and Vayatis, 2007). However, in our case, we do not want to emphasize the most common phrase over the rare but legal phrases. Therefore we use a simple pairwise criterion. Let $f(s, \theta)$ denote the score computed by our neural network with parameters θ , we minimize the ranking criterion:

$$\theta \mapsto \sum_{s \in \mathcal{S}} \sum_{w \in \mathcal{D}} \max \left\{ 0, 1 - f(s, \theta) + f(s^{(w)}, \theta) \right\}, \quad (18)$$

where \mathcal{S} is the set of text windows coming from our training corpus, \mathcal{D} is the dictionary of words, and $s^{(w)}$ denotes the text window obtained by replacing the central word of text window s by the word w .

Okanohara and Tsujii (2007) use a related approach to avoiding the entropy criteria using a binary classification approach (correct/incorrect phrase). Their work focuses on using a kernel classifier, and not on learning word embeddings as we do here.

4.3 Training Language Models

The language model network was trained by stochastic gradient minimization of the ranking criterion (18), sampling a sentence-word pair (s, w) at each iteration.

Since training times for such large scale systems are counted in weeks, it is not feasible to try many combinations of hyperparameters. It also makes sense to speedup the training time by initializing new networks with the embeddings computed by earlier networks. In particular, we found it expedient to train a succession of networks using increasingly large dictionaries, each network being initialized with the embeddings of the previous network. Successive dictionary sizes and switching times are chosen arbitrarily. (Bengio et al., 2009) provides a more detailed discussion of this the (as yet, poorly understood) “curriculum” process.

By analogy with biological cell lines, we have bred a few network lines. Within each line, child networks were initialized with the embeddings of their parents and trained on increasingly rich datasets with sometimes different parameters. Breeding decisions were made on the basis of the value of the ranking criterion (18) estimated on a validation set composed of one million words held out from the Wikipedia corpus.

9. However, Klein and Manning (2002) describe a rare example of realistic unsupervised grammar induction using a cross-entropy approach on binary-branching parsing trees, that is, by forcing the system to generate a hierarchical representation.

FRANCE	JESUS	XBOX	REDDISH	SCRATCHED	MEGABITS
454	1973	6909	11724	29869	87025
AUSTRIA	GOD	AMIGA	GREENISH	NAILED	OCTETS
BELGIUM	SATI	PLAYSTATION	BLUISH	SMASHED	MB/S
GERMANY	CHRIST	MSX	PINKISH	PUNCHED	BIT/S
ITALY	SATAN	IPOD	PURPLISH	POPPED	BAUD
GREECE	KALI	SEGA	BROWNISH	CRIMPED	CARATS
SWEDEN	INDRA	PSNUMBER	GREYISH	SCRAPED	KBIT/S
NORWAY	VISHNU	HD	GRAYISH	SCREWED	MEGAHERTZ
EUROPE	ANANDA	DREAMCAST	WHITISH	SECTIONED	MEGAPIXELS
HUNGARY	PARVATI	GEFORCE	SILVERY	SLASHED	GBIT/S
SWITZERLAND	GRACE	CAPCOM	YELLOWISH	RIPPED	AMPERES

Table 6: Word embeddings in the word lookup table of the language model neural network LM1 trained with a dictionary of size 100,000. For each column the queried word is followed by its index in the dictionary (higher means more rare) and its 10 nearest neighbors (using the Euclidean metric, which was chosen arbitrarily).

Very long training times make such strategies necessary for the foreseeable future: if we had been given computers ten times faster, we probably would have found uses for datasets ten times bigger. Of course this process makes very difficult to characterize the learning procedure. However we can characterize the final product.

In the following subsections, we report results obtained with two trained language models. The results achieved by these two models are representative of those achieved by networks trained on the full corpora.

- Language model LM1 has a window size $d_{win} = 11$ and a hidden layer with $n_{hu}^1 = 100$ units. The embedding layers were dimensioned like those of the supervised networks (Table 4). Model LM1 was trained on our first English corpus (Wikipedia) using successive dictionaries composed of the 5000, 10,000, 30,000, 50,000 and finally 100,000 most common WSJ words. The total training time was about four weeks.
- Language model LM2 has the same dimensions. It was initialized with the embeddings of LM1, and trained for an additional three weeks on our second English corpus (Wikipedia+Reuters) using a dictionary size of 130,000 words.

4.4 Embeddings

Both networks produce much more appealing word embeddings than in Section 3.4. Table 6 shows the ten nearest neighbours of a few randomly chosen query words for the LM1 model. The syntactic and semantic properties of the neighbours are clearly related to those of the query word. These results are far more satisfactory than those reported in Table 6 for embeddings obtained using purely supervised training of the benchmark NLP tasks.

Approach	POS (PWA)	CHUNK (F1)	NER (F1)	SRL (F1)
Benchmark Systems	97.24	94.29	89.31	77.92
NN+ITC	96.31	89.13	79.53	55.40
NN+STC	96.37	90.33	81.47	70.99
NN+ITC+LM1	97.05	91.91	85.68	58.18
NN+STC+LM1	97.10	93.65	87.58	73.84
NN+ITC+LM2	97.14	92.04	86.96	58.34
NN+STC+LM2	97.20	93.63	88.67	74.15

Table 7: Comparison in generalization performance of benchmark NLP systems with our neural network (NN) approach on POS, chunking, NER and SRL tasks. We report results with both the isolated tag criterion (ITC) and the sentence tag criterion (STC). We report with (LM n) performance of the networks trained from the language model embeddings (Table 6). Generalization performance is reported in per-word accuracy (PWA) for POS and F1 score for other tasks.

4.5 Semi-supervised Benchmark Results

Semi-supervised learning has been the object of much attention during the last few years (see Chapelle et al., 2006). Previous semi-supervised approaches for NLP can be roughly categorized as follows:

- Ad-hoc approaches such as (Rosenfeld and Feldman, 2007) for relation extraction.
- Self-training approaches, such as (Ueffing et al., 2007) for machine translation, and (McClosky et al., 2006) for parsing. These methods augment the labeled training set with examples from the unlabeled dataset using the labels predicted by the model itself. Transductive approaches, such as (Joachims, 1999) for text classification can be viewed as a refined form of self-training.
- Parameter sharing approaches such as (Ando and Zhang, 2005). Ando and Zhang propose a multi-task approach where they jointly train models sharing certain parameters. They train POS and NER models together with a language model consisting of predicting words given the surrounding tokens. This work should be seen as a linear counterpart of our approach. Using multilayer models vastly expands the parameter sharing opportunities (see Section 5).

Our approach simply consists of initializing the word lookup tables of the supervised networks with the embeddings computed by the language models. Supervised training is then performed as in section Section 3.4. In particular the supervised training stage is free to modify the lookup tables. This sequential approach is computationally convenient because it separates the lengthy training of the language models from the relatively fast training of the supervised networks. Once the language models are trained, we can perform multiple experiments on the supervised networks in a relatively short time.

Table 7 clearly shows that this simple initialization significantly boosts the generalization performance of the supervised networks for each task. It is worth mentioning the larger

language model led to even better performance. This suggests that we could still take advantage of even bigger unlabeled datasets.

4.6 Ranking and Language

There is a large agreement in the NLP community that syntax is a necessary prerequisite for semantic role labelling (Gildea and Palmer, 2002). This is why state-of-the-art semantic role labelling systems thoroughly exploit multiple parse trees. The parsers themselves (Charniak, 2000; Collins, 1999) contain considerable prior information about syntax.

Our system does not use such parse trees because we attempt to learn this information from the unlabelled data set. It is therefore legitimate to question whether our ranking criterion (18) has the conceptual capability to capture such a rich hierarchical information. At first glance, the ranking task appears unrelated to the induction of probabilistic grammars that underly standard parsing algorithms. The lack of hierarchical representation seems a fatal flaw (Chomsky, 1956).

However, ranking is closely related to a less known mathematical description of the language called *operator grammars* (Harris, 1968). Instead of directly studying the structure of a sentence, Harris studies the structure of the space of all sentences. Starting from a couple of elementary sentence forms, the space of all sentences is described by the successive application of sentence transformation operators. The sentence structure is revealed as a side effect of the successive transformations. Sentence transformations can also have a semantic interpretation.

Following the structural linguistic tradition, Harris describes procedures that leverage statistical properties of the language in order to discover the sentence transformation operators. Such procedures are obviously useful for machine learning approaches. In particular, he proposes a test to decide whether two sentences forms are or are not semantically related by a transformation operator. He first defines a ranking criterion similar to our ranking criterion (Harris, 1968, section 4.1):

“Starting for convenience with very short sentence forms, say ABC , we choose a particular word choice for all the classes, say B_qC_q , except one, in this case A ; for every pair of members A_i, A_j of that word class we ask how the sentence formed with one of the members, i.e. $A_iB_qC_q$ compares as to acceptability with the sentence formed with the other member, i.e. $A_jB_qC_q$.”

These *gradings* are then used to compare sentence forms:

“It now turns out that, given the graded n -tuples of words for a particular sentence form, we can find other sentences forms of the same word classes in which the same n -tuples of words produce the same grading of sentences.”

This is an indication that these two sentence forms exploit common words with the same syntactical function and possibly the same meaning. This observation forms the empirical basis for the construction of operator grammars that describe real-world natural languages such as English (Harris, 1982).

Therefore there are solid reasons to believe that the ranking criterion (18) has the conceptual potential to capture strong syntactical and semantical information. On the

other hand, the structure of our language models is probably too restrictive for such goals, and our current approach only exploits the word embeddings discovered during training.

5. Multi-Task Learning

It is generally accepted that features *trained* for one task can be useful for *related tasks*. This idea was already exploited in the previous section when certain language model features, namely the word embeddings, were used to initialize the supervised networks.

Multi-task learning (MTL) leverages this idea in a more systematic way. Models for all tasks of interests are *jointly trained* with an additional linkage between their trainable parameters in the hope of improving the generalization error. This linkage can take the form of a regularization term in the joint cost function that biases the models towards common representations. A much simpler approach consists in having the models *share certain parameters* defined a priori. Multi-task learning has a long history in machine learning and neural networks. Caruana (1997) gives a good overview of these past efforts.

5.1 Joint Decoding versus Joint Training

Multitask approaches do not necessarily involve joint training. For instance, modern speech recognition systems use Bayes rule to combine the outputs of an acoustic model trained on speech data and a language model trained on phonetic or textual corpora (Jelinek, 1976). This joint decoding approach has been successfully applied to structurally more complex NLP tasks. Sutton and McCallum (2005b) obtains improved results by combining the predictions of independently trained CRF models using a joint decoding process at test time that requires more sophisticated probabilistic inference techniques. On the other hand, Sutton and McCallum (2005a) obtain results somewhat below the state-of-the-art using joint decoding for SRL and syntactic parsing. Musillo and Merlo (2006) also describe a negative result at the same joint task.

Joint decoding invariably works by considering additional probabilistic dependency paths between the models. Therefore it defines an implicit supermodel that describes all the tasks in the same probabilistic framework. Separately training a submodel only makes sense when the training data blocks these additional dependency paths (in the sense of d-separation, Pearl, 1988). This implies that, without joint training, the additional dependency paths cannot directly involve unobserved variables. Therefore, the natural idea of discovering common internal representations across tasks requires joint training.

Joint training is relatively straightforward when the training sets for the individual tasks contain the same patterns with different labels. It is then sufficient to train a model that computes multiple outputs for each pattern (Suddarth and Holden, 1991). Using this scheme, Sutton et al. (2007) demonstrates improvements on POS tagging and noun-phrase chunking using jointly trained CRFs. However the joint labelling requirement is a limitation because such data is not often available. Miller et al. (2000) achieves performance improvements by jointly training NER, parsing, and relation extraction in a statistical parsing model. The joint labeling requirement problem was weakened using a predictor to fill in the missing annotations.

Ando and Zhang (2005) propose a setup that works around the joint labeling requirements. They define linear models of the form $f_i(x) = \mathbf{w}_i^\top \Phi(x) + \mathbf{v}_i^\top \Theta\Psi(x)$ where




Figure 3: Example of multitasking with NN. Task 1 and Task 2 are two tasks trained with the architecture presented in Figure 1. Lookup tables as well as the first hidden layer are shared. The last layer is task specific. The principle is the same with more than two tasks.

Approach	POS (PWA)	CHUNK (F1)	NER (F1)
Benchmark Systems	97.24	94.29	89.31
NN+STC+LM2	97.20	93.63	88.67
NN+STC+LM2+MTL	97.22	94.10	88.62

Table 8: Effect of multi-tasking on our neural architectures. We trained POS, CHUNK and NER in a MTL way. As a baseline, we show previous results of our system trained separately on each task. We also report benchmark systems performance for comparison.

f_i is the classifier for the i -th task with parameters \mathbf{w}_i and \mathbf{v}_i . Notations $\Phi(x)$ and $\Psi(x)$ represent engineered features for the pattern x . Matrix Θ maps the $\Psi(x)$ features into a low dimensional subspace common across all tasks. Each task is trained using its own examples without a joint labelling requirement. The learning procedure alternates the optimization of \mathbf{w}_i and \mathbf{v}_i for each task, and the optimization of Θ to minimize the average loss for all examples in all tasks. The authors also consider auxiliary unsupervised tasks for predicting substructures. They report excellent results on several tasks, including POS and NER.

5.2 Multi-Task Benchmark Results

Table 8 reports results obtained by jointly trained models for the POS, CHUNK and NER tasks using the same setup as Section 4.5. All three models share the lookup table parameters (2), and the first hidden layer parameters (4), as illustrated in Figure 3. Best results were obtained by enlarging the first hidden layer size to $n_{hu}^1 = 500$ in order to account for its shared responsibilities. The word embedding dimension was kept constant $d^0 = 50$ in order to reuse the language models of Section 4.5.

Training was achieved by minimizing the loss averaged across all tasks. This is easily achieved with stochastic gradient by alternatively picking examples for each task and applying (17) to all the parameters of the corresponding model, including the shared parameters. Note that this gives each task equal weight. Since each task uses the training sets described in Table 1, it is worth noticing that examples can come from quite different datasets. The generalization performance for each task was measured using the traditional testing data specified in Table 1. Fortunately, none of the training and test sets overlap across tasks.

Only the chunking task displays a significant performance improvement over the results reported in Section 4.5. The next section shows more directly that this improvement is related to the synergy between the POS and CHUNK task.

6. The Temptation

Results so far have been obtained by staying true to our *from scratch* philosophy. We have so far avoided specializing our architecture for any task, disregarding a lot of useful *a priori* NLP knowledge. We have shown that, thanks to very large unlabelled datasets, our generic neural networks can still achieve close to state-of-the-art performance by discovering useful features. This section explores what happens when we increase the level of engineering in our systems by incorporating some common techniques from the NLP literature. We often obtain further improvements. These figures are useful to quantify how far we went by leveraging large datasets instead of relying on a priori knowledge.

6.1 Stemming

Word suffixes in many western languages are strong predictors of the syntactic function of the word and therefore can benefit POS system. For instance, Ratnaparkhi (1996) uses inputs representing word suffixes and prefixes up to four characters. We achieve this in the POS task by adding discrete word features (Section 3.1.1) representing the last two characters of every word. The size of the suffix dictionary was 455. This led to a small improvement of the POS performance (Table 9, row NN+STC+LM2+Suffix2). We also tried suffixes obtained with the Porter (1980) stemmer and obtained the same performance as when using two character suffixes.

6.2 Gazetteers

State-of-the-art NER systems often use a large dictionary containing well known named entities (e.g. Florian et al., 2003). We restricted ourselves to the gazetteer provided by the CoNLL challenge, containing 8,000 locations, person names, organizations, and

Approach	POS (PWA)	CHUNK (F1)	NER (F1)	SRL
Benchmark Systems	97.24	94.29	89.31	77.92
NN+STC+LM2	97.20	93.63	88.67	74.15
NN+STC+LM2+Suffix2	97.29	—	—	—
NN+STC+LM2+Gazetteer	—	—	89.59	—
NN+STC+LM2+POS	—	94.32	88.67	—
NN+STC+LM2+CHUNK	—	—	—	74.72

Table 9: Comparison in generalization performance of benchmark NLP systems with our neural network (NN) using increasing task-specific engineering tricks. We report as a baseline the previous results obtained with network trained without task-specific tricks. The POS network was trained with two character word suffixes; the NER network was trained using the small CoNLL 2003 gazetteer; the CHUNK and NER networks were trained with additional POS features; and finally, the SRL network was trained with additional CHUNK features.

miscellaneous entities. We trained a NER network with 4 additional word features indicating whether the word is found in the gazetteer under one of these four categories. The resulting system displays a clear performance improvement (Table 9, row NN+STC+LM2+Gazetteer), establishing a new state-of-the-art NER performance.

6.3 Cascading

When one considers related tasks, it is reasonable to assume that tags obtained for one task can be useful for taking decisions in the other tasks. Conventional NLP systems often use features obtained from the output of other preexisting NLP systems. For instance, Shen and Sarkar (2005) describes a chunking system that uses POS tags as input; Florian et al. (2003) describes a NER system whose inputs include POS and CHUNK tags, as well as the output of two other NER classifiers. State-of-the-art SRL systems exploit parse trees (Gildea and Palmer, 2002; Punyakanok et al., 2005), related to CHUNK tags, and built using POS tags (Charniak, 2000; Collins, 1999).

Table 9 reports results obtained for the CHUNK and NER tasks by adding discrete word features (Section 3.1.1) representing the POS tags. In order to facilitate comparisons, instead of using the more accurate tags from our POS network, we use for each task the POS tags provided by the corresponding CoNLL challenge. We also report results obtained for the SRL task by adding word features representing the CHUNK tags. We consistently obtain moderate improvements.

6.4 Ensembles

Constructing ensembles of classifiers is a proven way to trade computational efficiency for generalization performance (Bell et al., 2007). Therefore it is not surprising that many NLP systems achieve state-of-the-art performance by combining the outputs of multiple classifiers. For instance, Kudo and Matsumoto (2001) use an ensemble of classifiers trained

Approach		POS (PWA)	CHUNK (F1)	NER (F1)
Benchmark Systems		97.24%	94.29%	89.31%
NN+STC+LM2+POS	worst	97.29%	93.99%	89.35%
NN+STC+LM2+POS	mean	97.31%	94.17%	89.65%
NN+STC+LM2+POS	best	97.35%	94.32%	89.86%
NN+STC+LM2+POS	ensemble	?	94.34%/94.41%	89.70%/89.88%

Table 10: Comparison in generalization performance for POS, CHUNK and NER tasks of the networks obtained using by combining ten training runs with different initialization.

with different tagging conventions (see Section 3.2.3). Winning a challenge is of course a legitimate objective. Yet it is often difficult to figure out which ideas are most responsible for the state-of-the-art performance of a large ensemble.

Because neural networks are nonconvex, training runs with different initial parameters usually give different solutions. Table 10 reports results obtained for the CHUNK and NER task after ten training runs with random initial parameters. Simply averaging the predictions of the 10 resulting networks leads to a small improvement over the best network. This improvement comes of course at the expense of a ten fold increase of the running time. On the other hand, multiple training times could be improved using smart sampling strategies (Neal, 1996).

We can also observe that the performance variability is not very large. The local minima found by the training algorithm are usually good local minima, thanks to the oversized parameter space and to the noise induced by the stochastic gradient procedure (LeCun et al., 1998). In order to reduce the variance in our experimental results, we always use the same initial parameters for networks trained on the same task (except of course for the results reported in Table 10.)

6.5 Parsing

Gildea and Palmer (2002) offers several arguments suggesting that syntactical parsing is a necessary prerequisite for the SRL task. The CoNLL 2005 SRL benchmark task provides parse trees computed using *both* the Charniak (2000) and Collins (1999) parsers. State-of-the-art systems often exploit additional parse trees such as the k top ranking parse trees (Koomen et al., 2005; Haghighi et al., 2005).

In contrast our SRL networks so far do not use parse trees at all. They rely instead on internal representations transferred from a language model trained with an objective function that captures a lot of syntactical information (see Section 4.6). It is therefore legitimate to question whether this approach is an acceptable lightweight replacement for parse trees.

We answer this question by providing parse tree information as input to our system. We have limited ourselves to the Charniak parse tree provided with the CoNLL 2005 data. The parse tree information is represented using additional word features encoding the parse tree information at a given depth. We call level 0 the information associated with the leaves

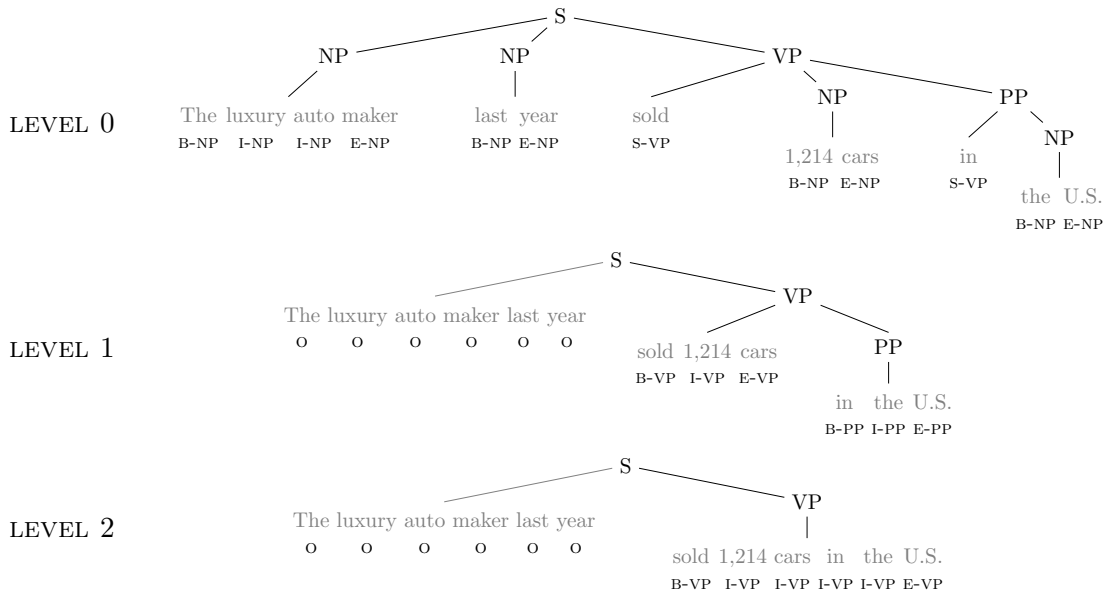


Figure 4: Charniak parse tree for the sentence “*The luxury auto maker last year sold 1,214 cars in the U.S.*”. Level 0 is the original tree. Levels 1 to 4 are obtained by successively collapsing terminal tree branches. For each level, words receive tags describing the segment associated with the corresponding leaf. All words receive tag “O” at level 3 in this example.

of the parse tree. We obtain levels 1 to 4 by repeatedly trimming the leaves as shown in Figure 4. Segment boundaries and node labels are encoded using IOBES word tags as discussed in Sections 3.2.3 and 3.1.1.

Experiments were performed using the LM1 language model using the same network architectures (see Table 4) and using additional word features of dimension 5 for each parse tree levels. Table 11 reports the performance improvements obtained by providing increasing levels of parse tree information. Level 0 alone increases the F1 score by almost 1.5%. Additional levels yields diminishing returns. The top performance reaches 76.06% F1 score. This compares nicely with the state-of-the-art system which uses *six parse trees* instead of one. Koomen et al. (2005) also report a 74.76% F1 score on the validation set using only the Charniak parse tree.

6.6 SENNA: Engineering a Sweet Spot

It is surprising to observe that adding chunking features into the semantic role labelling network (Table 9) performs significantly worse than adding features describing the level 0 of the Charniak parse tree (Table 11). It appears that the parse trees identify leaf sentence segments that are often smaller than those identified by the chunking tags. However we can easily train a second chunking network to identify the segments delimited by the leaves of the Penn Treebank parse trees. We can then retrain the semantic role labelling network using this information as input. Table 12 summarizes the results achieved in this manner.

Approach	SRL (test set F1)
Benchmark System (six parse trees)	77.92
Benchmark System (top Charniak parse tree only, valid. set)	74.76
NN+STC+LM2	74.15
NN+STC+LM2+Charniak (level 0 only)	75.65
NN+STC+LM2+Charniak (levels 0 & 1)	75.81
NN+STC+LM2+Charniak (levels 0 to 2)	76.05
NN+STC+LM2+Charniak (levels 0 to 3)	75.89
NN+STC+LM2+Charniak (levels 0 to 4)	76.06

Table 11: Generalization performance on the SRL task of our NN architecture compared with the benchmark system. We show performance of our system fed with different levels of depth of the Charniak parse tree. We report previous results of our architecture with no parse tree as a baseline. The benchmark system (Koomen et al., 2005) uses six parse trees. However they also report the performance obtained on the validation set using only the top Charniak parse tree. We report this result because it is more comparable to ours.

Task		Benchmark	SENNA
Part of Speech (POS)	(Accuracy)	97.24 %	97.29 %
Chunking (CHUNK)	(F1)	94.29 %	94.32 %
Named Entity Recognition (NER)	(F1)	89.31 %	89.59 %
Parse Tree level 0 (PT0)	(F1)	91.94 %	92.25 %
Semantic Role Labelling (SRL)	(F1)	77.92 %	75.55 %

Table 12: Performance of the engineered sweet spot (SENNA) on various tagging tasks. The PT0 task replicates the sentence segmentation of the parse tree leaves. The corresponding benchmark score measures the quality of the Charniak parse tree leaves relative to the Penn Treebank gold parse trees.

The resulting system is fast because it uses a simple architecture, self-contained, because it does not rely on the output of existing NLP system, and accurate because it offers state-of-the-art or near state-of-the-art performance on a variety of tagging tasks.

POS System	Time (s.)	SRL System	Time (s.)
Toutanova et al. (2003)	1065	Koomen et al. (2005)	6253
Shen et al. (2007)	833		
SENNA	4	SENNA	51

Table 13: Runtime speed comparison between state-of-the-art systems and our approach (SENNA). We give the runtime in seconds for running both the POS and SRL taggers on their respective testing sets.

The runtime version¹⁰ contains about 2000 lines of C code, runs in less than 100MB of memory, and needs only a couple milliseconds per word to compute all the tags. Table 13 compares the tagging speeds for our system and for the few available state-of-the-art systems: the Toutanova et al. (2003) and Shen et al. (2007) POS taggers¹¹, and the Koomen et al. (2005) SRL system.¹² All programs were run on a single 3GHz Intel core. The POS taggers were run with Sun Java 1.6 with a large enough memory allocation to reach their top tagging speed. The beam size of the Shen tagger was set to 3 as recommended in the paper. Regardless of implementation differences, it is clear that our neural networks run considerably faster. They also require much less memory. Our POS tagger runs in 32MB of RAM. The Shen and Toutanova taggers slow down very significantly when the Java machine is given less than 2.2GB and 1.1GB of RAM respectively.

7. Critical Discussion

Although we believe that this contribution represents a step towards the “NLP from scratch” objective, we are keenly aware that both our goal and our means can be criticized.

The main criticism of our goal can be summarized as follows. Over the years, the NLP community has developed a considerable expertise in engineering effective NLP features. Why should they forget this painfully acquired expertise and instead painfully acquire the skills required to train large neural networks? As mentioned in our introduction, we observe that no single NLP task really covers the goals of NLP. Therefore we believe that excessive task-specific engineering is not desirable. But we also recognize how much our neural networks owe to previous NLP task-specific research.

The main criticism of our means is easier to address. Why did we choose to rely on a twenty year old technology, namely multilayer neural networks? We were simply attracted by their ability to discover hidden representations using a stochastic learning algorithm that scales linearly with the number of examples. Most of the neural network technology necessary for our work has been described ten years ago (e.g. Le Cun et al., 1998). However, if we had decided ten years ago to train the language model network LM2 using a vintage computer, training would only be nearing completion today. Training algorithms that scale linearly are most able to benefit from such tremendous progress in computer hardware.

8. Conclusion

We have presented a deep neural network architecture that can handle a number of NLP tasks with both speed and accuracy. The design of this system was determined by our desire to avoid task-specific engineering as much as possible. Instead we rely on very large unlabeled datasets and let the training algorithm discover internal representations that prove useful for all the tasks of interest. Using this strong basis, we have engineered a fast and efficient “all purpose” NLP tagger that we hope will prove useful to the community.

10. Available at <http://ml.nec-labs.com/senna>.

11. Available at <http://nlp.stanford.edu/software/tagger.shtml>.

12. Available at <http://l2r.cs.uiuc.edu/~cogcomp/asoftware.php?skey=SRL>.

Acknowledgments

We acknowledge the persistent support of NEC for this research effort. We thank Yoshua Bengio, Samy Bengio, Eric Cosatto, Vincent Etter, Hans-Peter Graf, Ralph Grishman, and Vladimir Vapnik for their useful feedback and comments.

References

- R. K. Ando and T. Zhang. A framework for learning predictive structures from multiple tasks and unlabeled data. *JMLR*, 6:1817–1953, 11 2005.
- R. M. Bell, Y. Koren, and C. Volinsky. The BellKor solution to the Netflix Prize. Technical report, AT&T Labs, 2007. <http://www.research.att.com/~volinsky/netflix>.
- Y. Bengio and R. Ducharme. A neural probabilistic language model. In *NIPS 13*, 2001.
- Y. Bengio, J. Louradour, R. Collobert, and J. Weston. Curriculum learning. In *International Conference on Machine Learning, ICML*, 2009.
- L. Bottou. Stochastic gradient learning in neural networks. In *Proceedings of Neuro-Nîmes 91*, Nîmes, France, 1991. EC2.
- L. Bottou and P. Gallinari. A framework for the cooperation of learning algorithms. In D. Touretzky and R. Lippmann, editors, *Advances in Neural Information Processing Systems*, volume 3. Morgan Kaufmann, Denver, 1991.
- L. Bottou, Y. LeCun, and Yoshua Bengio. Global training of document processing systems using graph transformer networks. In *Proc. of Computer Vision and Pattern Recognition*, pages 489–493, Puerto-Rico, 1997. IEEE.
- J. S. Bridle. Probabilistic interpretation of feedforward classification network outputs, with relationships to statistical pattern recognition. In F. Fogelman Soulié and J. Hérault, editors, *Neurocomputing: Algorithms, Architectures and Applications*, pages 227–236. NATO ASI Series, 1990.
- P. F. Brown, V. J. Della Pietra, R. L. Mercer, S. A. Della Pietra, and J. C. Lai. An estimate of an upper bound for the entropy of english. *Computational Linguistics*, 18(1):31–41, 1992.
- C. J. C. Burges, R. Ragno, and Quoc Viet Le. Learning to rank with nonsmooth cost functions. In B. Schölkopf, J. Platt, and T. Hoffman, editors, *Advances in Neural Information Processing Systems 19*, pages 193–200. MIT Press, Cambridge, MA, 2007.
- R. Caruana. Multitask Learning. *Machine Learning*, 28(1):41–75, 1997.
- O. Chapelle, B. Schölkopf, and A. Zien. *Semi-Supervised Learning*. Adaptive computation and machine learning. MIT Press, Cambridge, Mass., USA, 09 2006.

- E. Charniak. A maximum-entropy-inspired parser. *Proceedings of the first conference on North American chapter of the Association for Computational Linguistics*, pages 132–139, 2000.
- Hai Leong Chieu. Named entity recognition with a maximum entropy approach. In *In Proceedings of the Seventh Conference on Natural Language Learning (CoNLL-2003)*, pages 160–163, 2003.
- N. Chomsky. Three models for the description of language. *IRE Transactions on Information Theory*, 2(3):113–124, September 1956.
- S. Cl  men  on and N. Vayatis. Ranking the best instances. *Journal of Machine Learning Research*, 8:2671–2699, December 2007.
- W. W. Cohen, R. E. Schapire, and Y. Singer. Learning to order things. *Journal of Artificial Intelligence Research*, 10:243–270, 1998.
- T. Cohn and P. Blunsom. Semantic role labelling with tree conditional random fields. In *Ninth Conference on Computational Natural Language (CoNLL)*, 2005.
- M. Collins. *Head-Driven Statistical Models for Natural Language Parsing*. PhD thesis, University of Pennsylvania, 1999.
- T. Cover and R. King. A convergent gambling estimate of the entropy of english. *IEEE Transactions on Information Theory*, 24(4):413–421, July 1978.
- R. Florian, A. Ittycheriah, H. Jing, and T. Zhang. Named entity recognition through classifier combination. In *Proceedings of the seventh conference on Natural language learning at HLT-NAACL 2003*, pages 168–171. Association for Computational Linguistics, 2003.
- D. Gildea and M. Palmer. The necessity of parsing for predicate argument recognition. *Proceedings of the 40th Annual Meeting of the ACL*, pages 239–246, 2002.
- J. Gim  nez and L. M  rquez. SVMTool: A general POS tagger generator based on support vector machines. In *Proceedings of the 4th International Conference on Language Resources and Evaluation (LREC’04)*, 2004.
- A. Haghighi, K. Toutanova, and C. D. Manning. A joint model for semantic role labeling. In *Proceedings of the Ninth Conference on Computational Natural Language Learning (CoNLL-2005)*. Association for Computational Linguistics, June 2005.
- Z. S. Harris. *Mathematical Structures of Language*. John Wiley & Sons Inc., 1968.
- Z. S. Harris. *A Grammar of English on Mathematical Principles*. John Wiley & Sons Inc., 1982.
- D. Heckerman, D. M. Chickering, C. Meek, R. Rounthwaite, and C. Kadie. Dependency networks for inference, collaborative filtering, and data visualization. *Journal of Machine Learning Research*, 1:49–75, 2001. ISSN 1532-4435.

- F. Jelinek. Continuous speech recognition by statistical methods. *Proceedings of the IEEE*, 64(4):532–556, 1976.
- T. Joachims. Transductive inference for text classification using support vector machines. In *ICML*, 1999.
- D. Klein and C. D. Manning. Natural language grammar induction using a constituent-context model. In Thomas G. Dietterich, Suzanna Becker, and Zoubin Ghahramani, editors, *Advances in Neural Information Processing Systems 14*, pages 35–42. MIT Press, Cambridge, MA, 2002.
- P. Koomen, V. Punyakanok, D. Roth, and W. Yih. Generalized inference with multiple semantic role labeling systems (shared task paper). In Ido Dagan and Dan Gildea, editors, *Proc. of the Annual Conference on Computational Natural Language Learning (CoNLL)*, pages 181–184, 2005.
- T. Kudo and Y. Matsumoto. Chunking with support vector machines. In *In Proceedings of the 2nd Meeting of the North American Association for Computational Linguistics: NAACL 2001*, pages 1–8. Association for Computational Linguistics, 2001.
- T. Kudoh and Y. Matsumoto. Use of support vector learning for chunk identification. In *Proceedings of CoNLL-2000 and LLL-2000*, pages 142–144, 2000.
- J. Lafferty, A. McCallum, and F. Pereira. Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *Eighteenth International Conference on Machine Learning, ICML*, 2001.
- Y. Le Cun, L. Bottou, Y. Bengio, and P. Haffner. Gradient based learning applied to document recognition. *Proceedings of IEEE*, 86(11):2278–2324, 1998.
- Y. LeCun. A learning scheme for asymmetric threshold networks. In *Proceedings of Cognitiva 85*, pages 599–604, Paris, France, 1985.
- Y. LeCun, L. Bottou, G. B. Orr, and K.-R. Müller. Efficient backprop. In G.B. Orr and K.-R. Müller, editors, *Neural Networks: Tricks of the Trade*, pages 9–50. Springer, 1998.
- D. D. Lewis, Y. Yang, T. G. Rose, and F. Li. Rcv1: A new benchmark collection for text categorization research. *Journal of Machine Learning Research*, 5:361–397, 2004.
- N. Littlestone. Learning quickly when irrelevant attributes abound: A new linear-threshold algorithm. In *Machine Learning*, pages 285–318, 1988.
- A. McCallum and Wei Li. Early results for named entity recognition with conditional random fields, feature induction and web-enhanced lexicons. In *Proceedings of the seventh conference on Natural language learning at HLT-NAACL 2003*, pages 188–191. Association for Computational Linguistics, 2003.
- D. McClosky, E. Charniak, and M. Johnson. Effective self-training for parsing. *Proceedings of HLT-NAACL 2006*, 2006.

- R. McDonald, K. Crammer, and F. Pereira. Flexible text segmentation with structured multilabel classification. In *HLT '05: Proceedings of the conference on Human Language Technology and Empirical Methods in Natural Language Processing*, pages 987–994. Association for Computational Linguistics, 2005.
- S. Miller, H. Fox, L. Ramshaw, and R. Weischedel. A novel use of statistical parsing to extract information from text. *6th Applied Natural Language Processing Conference*, 2000.
- G. Musillo and P. Merlo. Robust Parsing of the Proposition Bank. *ROMAND 2006: Robust Methods in Analysis of Natural language Data*, 2006.
- R. M. Neal. *Bayesian Learning for Neural Networks*. Number 118 in Lecture Notes in Statistics. Springer-Verlag, New York, 1996.
- D. Okanohara and J. Tsujii. A discriminative language model with pseudo-negative samples. *Proceedings of the 45th Annual Meeting of the ACL*, pages 73–80, 2007.
- M. Palmer, D. Gildea, and P. Kingsbury. The proposition bank: An annotated corpus of semantic roles. *Comput. Linguist.*, 31(1):71–106, 2005. ISSN 0891-2017. doi: <http://dx.doi.org/10.1162/0891201053630264>.
- J. Pearl. *Probabilistic Reasoning in Intelligent Systems*. Morgan Kaufman, San Mateo, 1988.
- D. C. Plaut and G. E. Hinton. Learning sets of filters using back-propagation. *Computer Speech and Language*, 2:35–61, 1987.
- M. F. Porter. An algorithm for suffix stripping. *Program*, 14(3):130–137, 1980.
- S. Pradhan, K. Hacioglu, W. Ward, J. H. Martin, and D. Jurafsky. Semantic role chunking combining complementary syntactic views. In *Proceedings of the Ninth Conference on Computational Natural Language Learning (CoNLL-2005)*, pages 217–220. Association for Computational Linguistics, June 2005.
- V. Punyakanok, D. Roth, and W. Yih. The necessity of syntactic parsing for semantic role labeling. In *IJCAI*, pages 1117–1123, 2005.
- A. Ratnaparkhi. A maximum entropy model for part-of-speech tagging. In Eric Brill and Kenneth Church, editors, *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, pages 133–142. Association for Computational Linguistics, 1996.
- B. Rosenfeld and R. Feldman. Using Corpus Statistics on Entities to Improve Semi-supervised Relation Extraction from the Web. *Proceedings of the 45th Annual Meeting of the ACL*, pages 600–607, 2007.
- D. E. Rumelhart, G. E. Hinton, and R. J. Williams. Learning internal representations by back-propagating errors. In D.E. Rumelhart and J. L. McClelland, editors, *Parallel Distributed Processing: Explorations in the Microstructure of Cognition*, volume 1, pages 318–362. MIT Press, 1986.

- H. Schwenk and J. L. Gauvain. Connectionist language modeling for large vocabulary continuous speech recognition. In *IEEE International Conference on Acoustics, Speech, and Signal Processing*, pages 765–768, 2002.
- Fei Sha and F. Pereira. Shallow parsing with conditional random fields. In *NAACL '03: Proceedings of the 2003 Conference of the North American Chapter of the Association for Computational Linguistics on Human Language Technology*, pages 134–141. Association for Computational Linguistics, 2003.
- C. E. Shannon. Prediction and entropy of printed english. *Bell Systems Technical Journal*, 30:50–64, 1951.
- H. Shen and A. Sarkar. Voting between multiple data representations for text chunking. *Advances in Artificial Intelligence*, pages 389–400, 2005.
- L. Shen, G. Satta, and A. K. Joshi. Guided learning for bidirectional sequence classification. In *Proceedings of the 45th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2007.
- S. C. Suddarth and A. D. C. Holden. Symbolic-neural systems and the use of hints for developing complex systems. *International Journal of Man-Machine Studies*, 35(3):291–311, 1991.
- X. Sun, L.-P. Morency, D. Okanohara, and J. Tsujii. Modeling latent-dynamic in shallow parsing: a latent conditional model with improved inference. In *COLING '08: Proceedings of the 22nd International Conference on Computational Linguistics*, pages 841–848. Association for Computational Linguistics, 2008.
- C. Sutton and A. McCallum. Joint parsing and semantic role labeling. In *Proceedings of CoNLL-2005*, pages 225–228, 2005a.
- C. Sutton and A. McCallum. Composition of conditional random fields for transfer learning. *Proceedings of the conference on Human Language Technology and Empirical Methods in Natural Language Processing*, pages 748–754, 2005b.
- C. Sutton, A. McCallum, and K. Rohanimanesh. Dynamic Conditional Random Fields: Factorized Probabilistic Models for Labeling and Segmenting Sequence Data. *JMLR*, 8: 693–723, 2007.
- W. J. Teahan and J. G. Cleary. The entropy of english using ppm-based models. In *In Data Compression Conference (DCC'96)*, pages 53–62. IEEE Computer Society Press, 1996.
- K. Toutanova, D. Klein, C. D. Manning, and Y. Singer. Feature-rich part-of-speech tagging with a cyclic dependency network. In *HLT-NAACL*, 2003.
- N. Ueffing, G. Haffari, and A. Sarkar. Transductive learning for statistical machine translation. *Proceedings of the 45th Annual Meeting of the ACL*, pages 25–32, 2007.
- A. Waibel, T. Hanazawa abd G. Hinton, K. Shikano, and K.J. Lang. Phoneme recognition using time-delay neural networks. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 37(3):328–339, 1989.